

# Q1 Write-Up

## Linux Scheduler

### Working

The following changes were made for this implementation:

In the sched.h file present in include/linux/sched.h, I introduced a new u64 type as delaytime in the sched\_entity struct, representing the delay that we need to pass. In the core.c file present in kernel/sched/core.c, in the \_\_sched\_fork() function, I added a new attribute to the task\_struct p as p->se.delaytime and initialized it to 0. After that in the scheduler file fair.c present in kernel/sched/fair.c, in the update\_curr() function, after the vruntime is updated by the calc\_delta\_fair() amount, I further updated the vruntime by the delaytime we initialized by doing curr->vruntime += curr->delaytime; Hence we modified the CFS scheduler so that the vruntime is updated according to the delay passed.

To define the system call I modified the sys.c file present in kernel/sys.c, and used the SYSCALL\_DEFINE2 macro that has parameters as the PID of the process whose vruntime should be changed and the delta amount time (in milli-seconds) by which the selection should be delayed. To get the task\_struct of the pid, first I used the find\_get\_pid() function which returns a pid struct and used pid\_task() to get the task struct of this pid. After that, I have updated the delaytime of our task\_struct variable 'p' by increasing the delaytime by the delta (multiplied by 10^6 to get in nanoseconds) using the statement p->se.delaytime += delta\*(1000000); Now, whenever the vruntime is calculated for the process, it will also get updated by this delta delay (as we defined in fair.c).

Finally, I modified the syscall\_64.tbl file present in arch/x86/entry/syscalls and added the new syscall name as 'change\_vruntime' and number as 449.